

Continuous All k Nearest Neighbor Queries in Smartphone Networks- **Proximity Sensing, Part I**

Source: <https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=6341377>

Outline

- **Motivation**
- System Model and Problem Formulation
- Proximity Algorithm
- Experimental Evaluation

Source: <https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=6341377>

Smartphones

- **Smartphone: A powerful sensing device!**
 - **Processing:** 2x2.73 GHz Octa-core Samsung Galaxy S10)
 - **RAM & Flash Storage:** 6GB & 128GB, respectively
 - **Networking:** WiFi, 3G (Mbps) / 4G (100Mbps–1Gbps)
 - **Sensing:** Proximity, Ambient Light, Accelerometer, Microphone, Geographic Coordinates based on AGPS (fine), WiFi or Cellular Towers (coarse) & etc.
- ***In-House Applications!***



SmartTrace
(ICDE'09, MDM'09, TKDE'12)



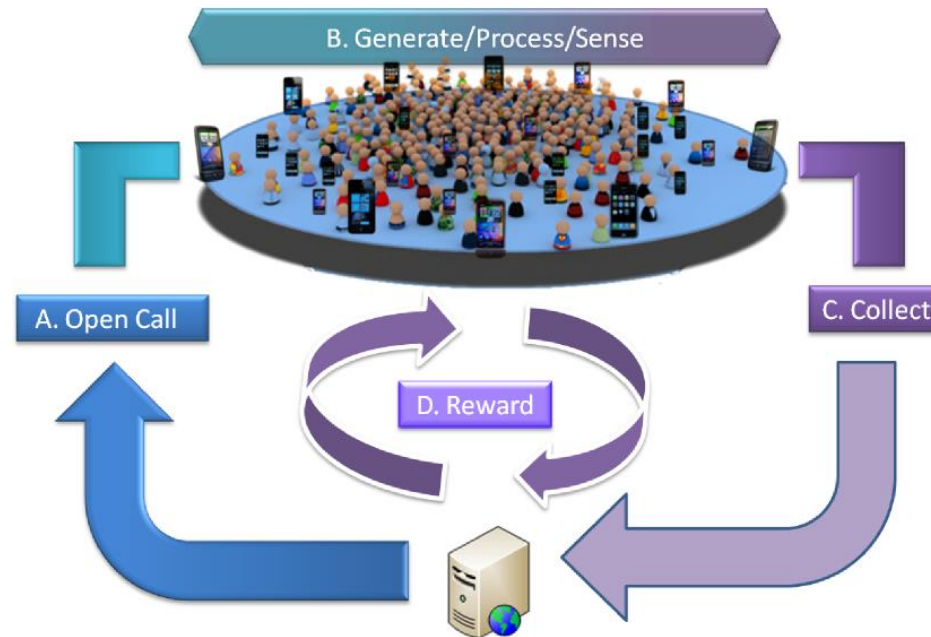
SmartP2P
(MDM'11, MDM'12)



Airplace
(MobiSys'12, MDM'12)

Motivation

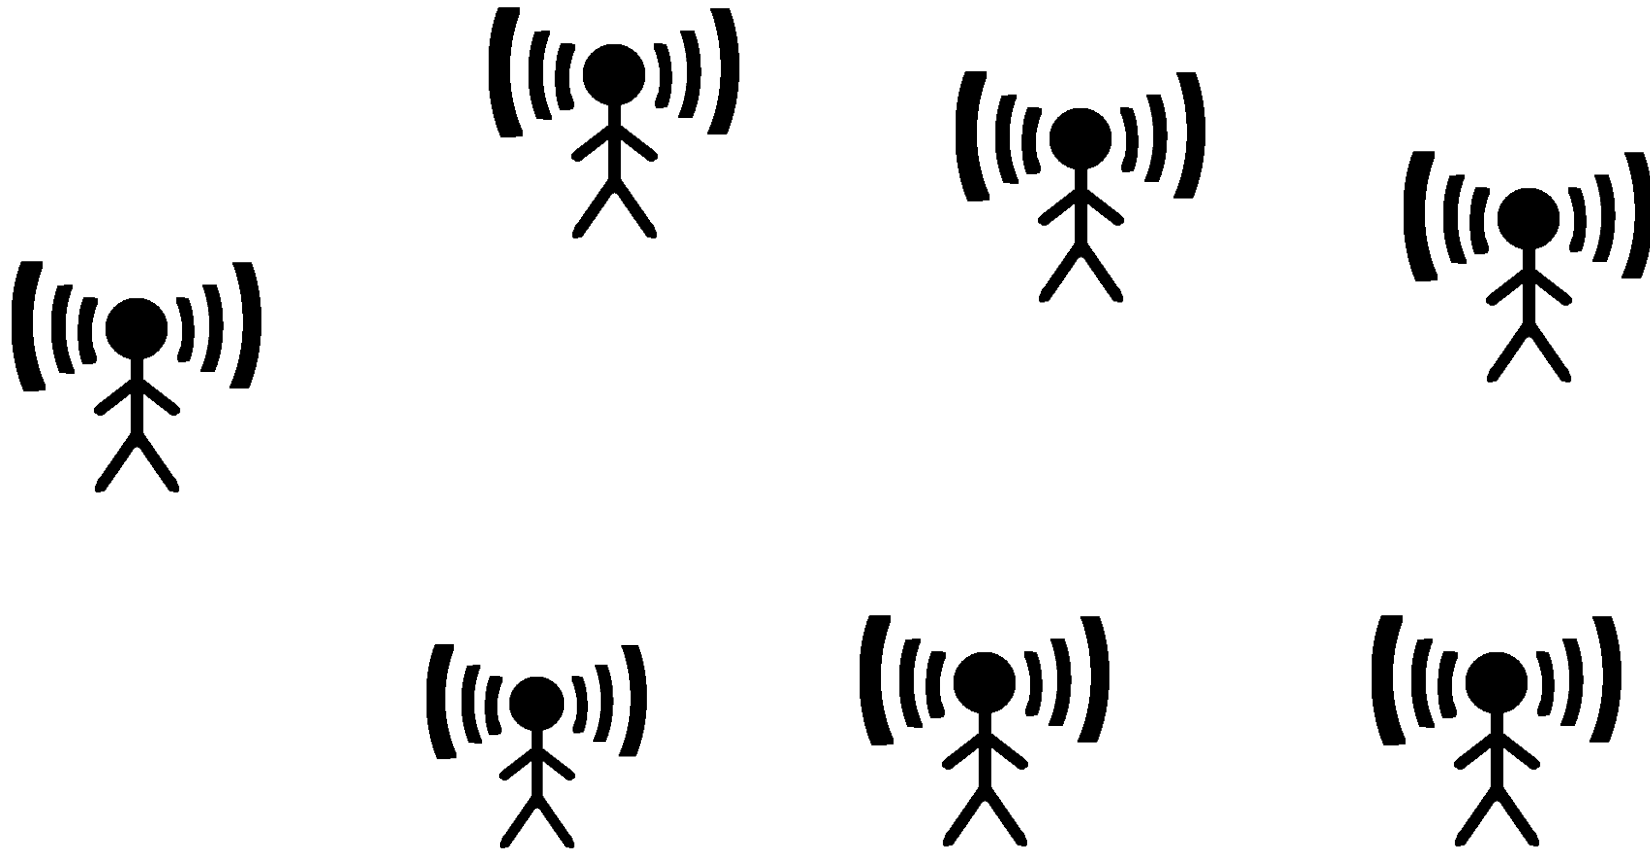
- Crowdsourcing with Smartphones
 - A smartphone crowd is constantly moving and sensing providing large amounts of opportunistic data that enables new services and applications



Source: <https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=6341377>

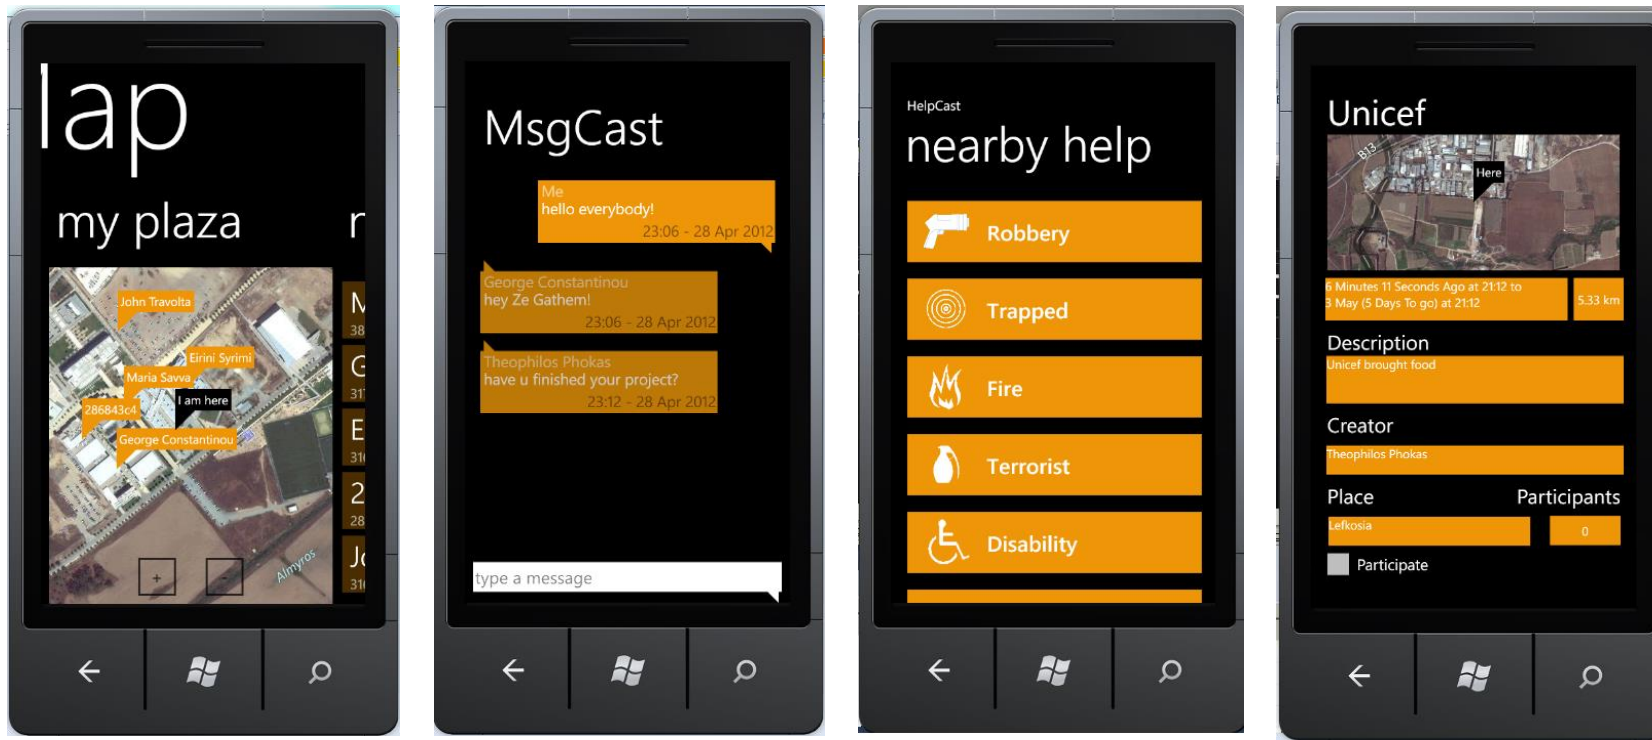
Continuous k Nearest Neighbor Queries

Find 2 Closest Neighbors for 1 User



Motivation

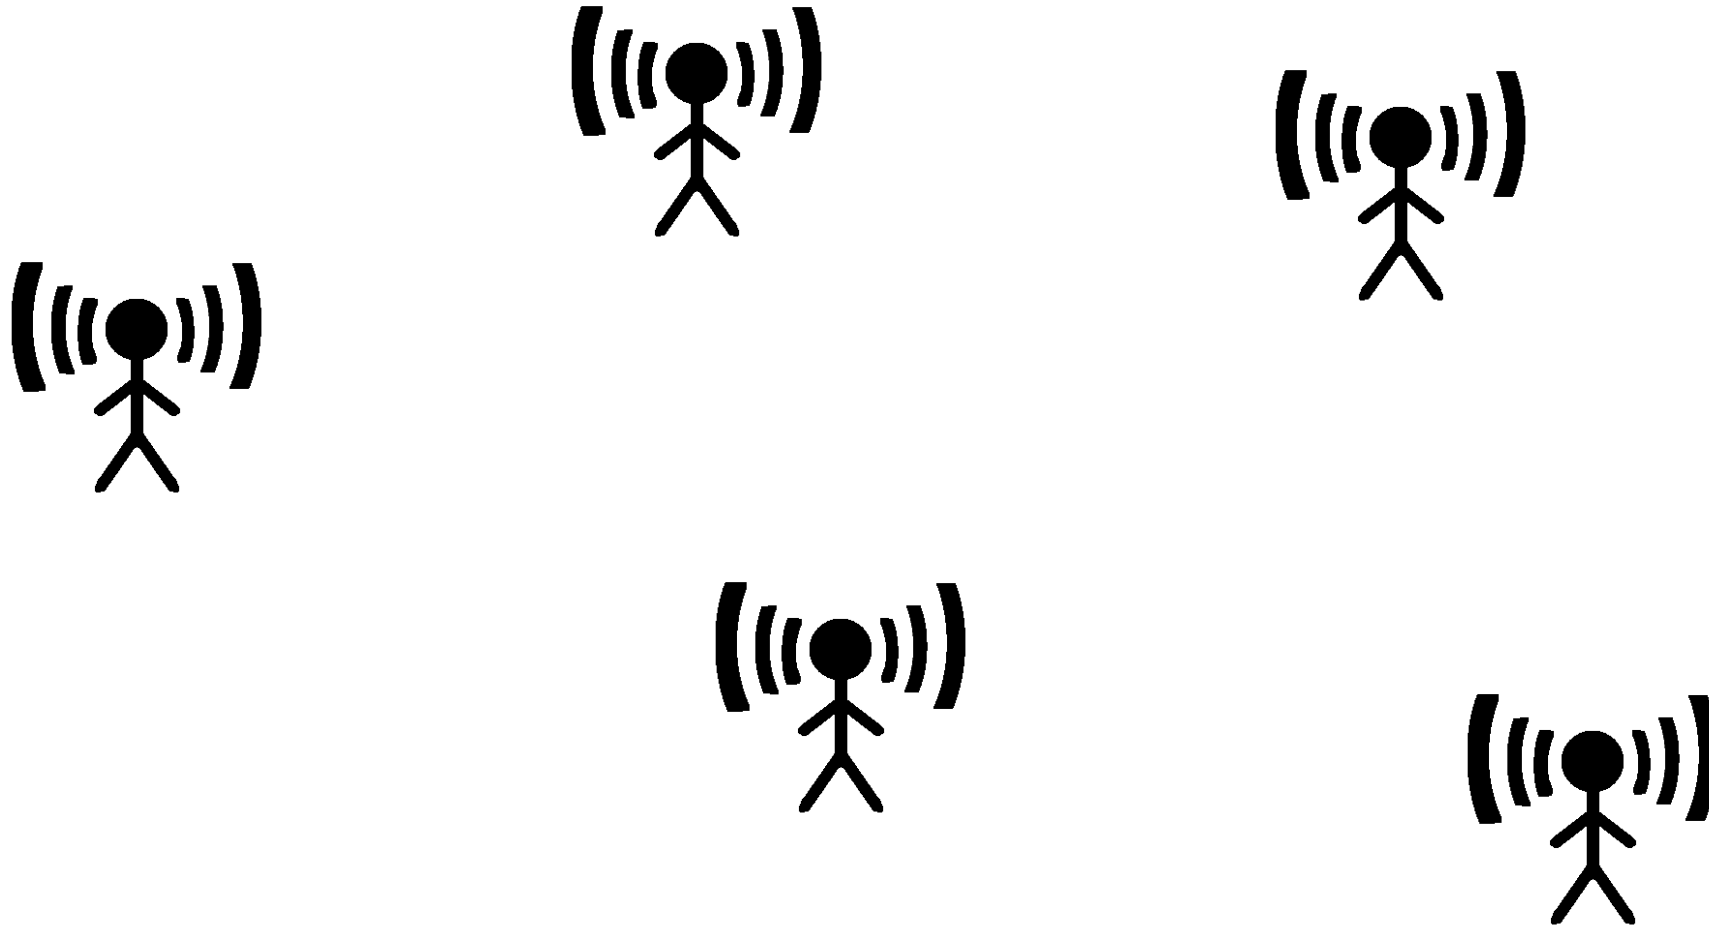
“Create a Framework for Efficient Proximity Interactions”



Screenshots from a prototype system for Windows Phone
<http://www.zegathem.com/>

Continuous All K Nearest Neighbor (CAkNN) Queries

Find 2 Closest Neighbors for ALL User

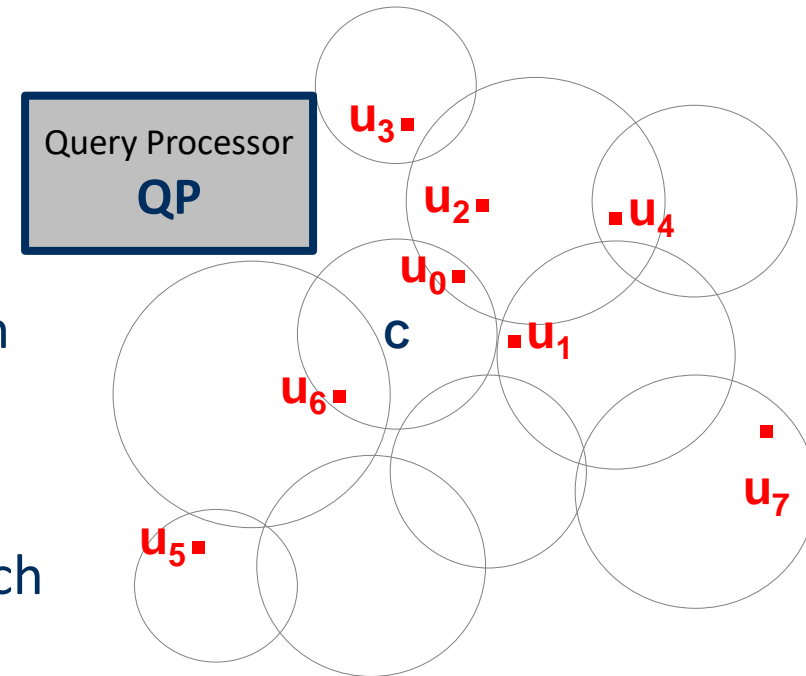


Outline

- Motivation
- **System Model and Problem Formulation**
- Proximity Algorithm
- Experimental Evaluation

System Model

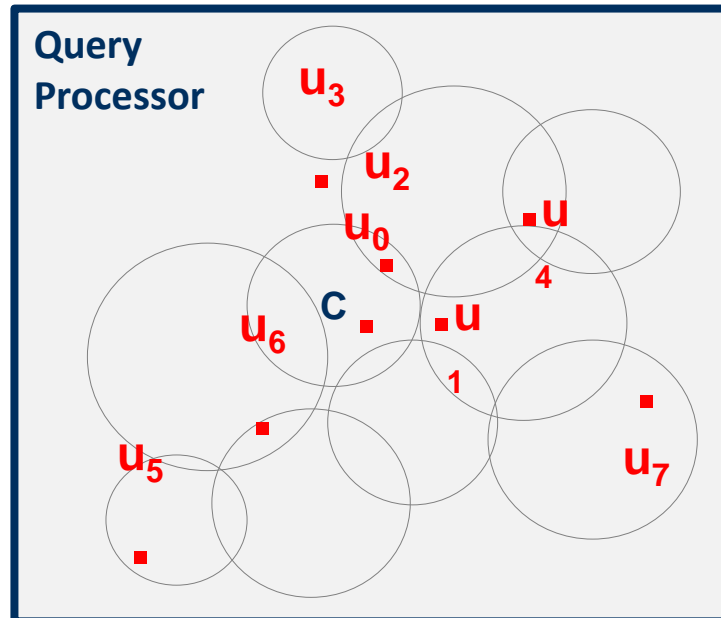
- A set of users moving in the plane of a region ($u_1, u_2, \dots u_n$)
- Area covered by a set of **Network Connectivity Points (NCP)**
 - Each *NCP* creates the notion of a *cell*
 - W.l.o.g., let the cell be represented by a circular area with an arbitrary radius
- A mobile user **u** is **serviced** at any given time point by one *NCP*.
- There is some **centralized service**, denoted as **QP (Query Processor)**, which is aware of the coverage of each NCP.
- Each user **u** reports its **positional information** to **QP** regularly



Problem Definition

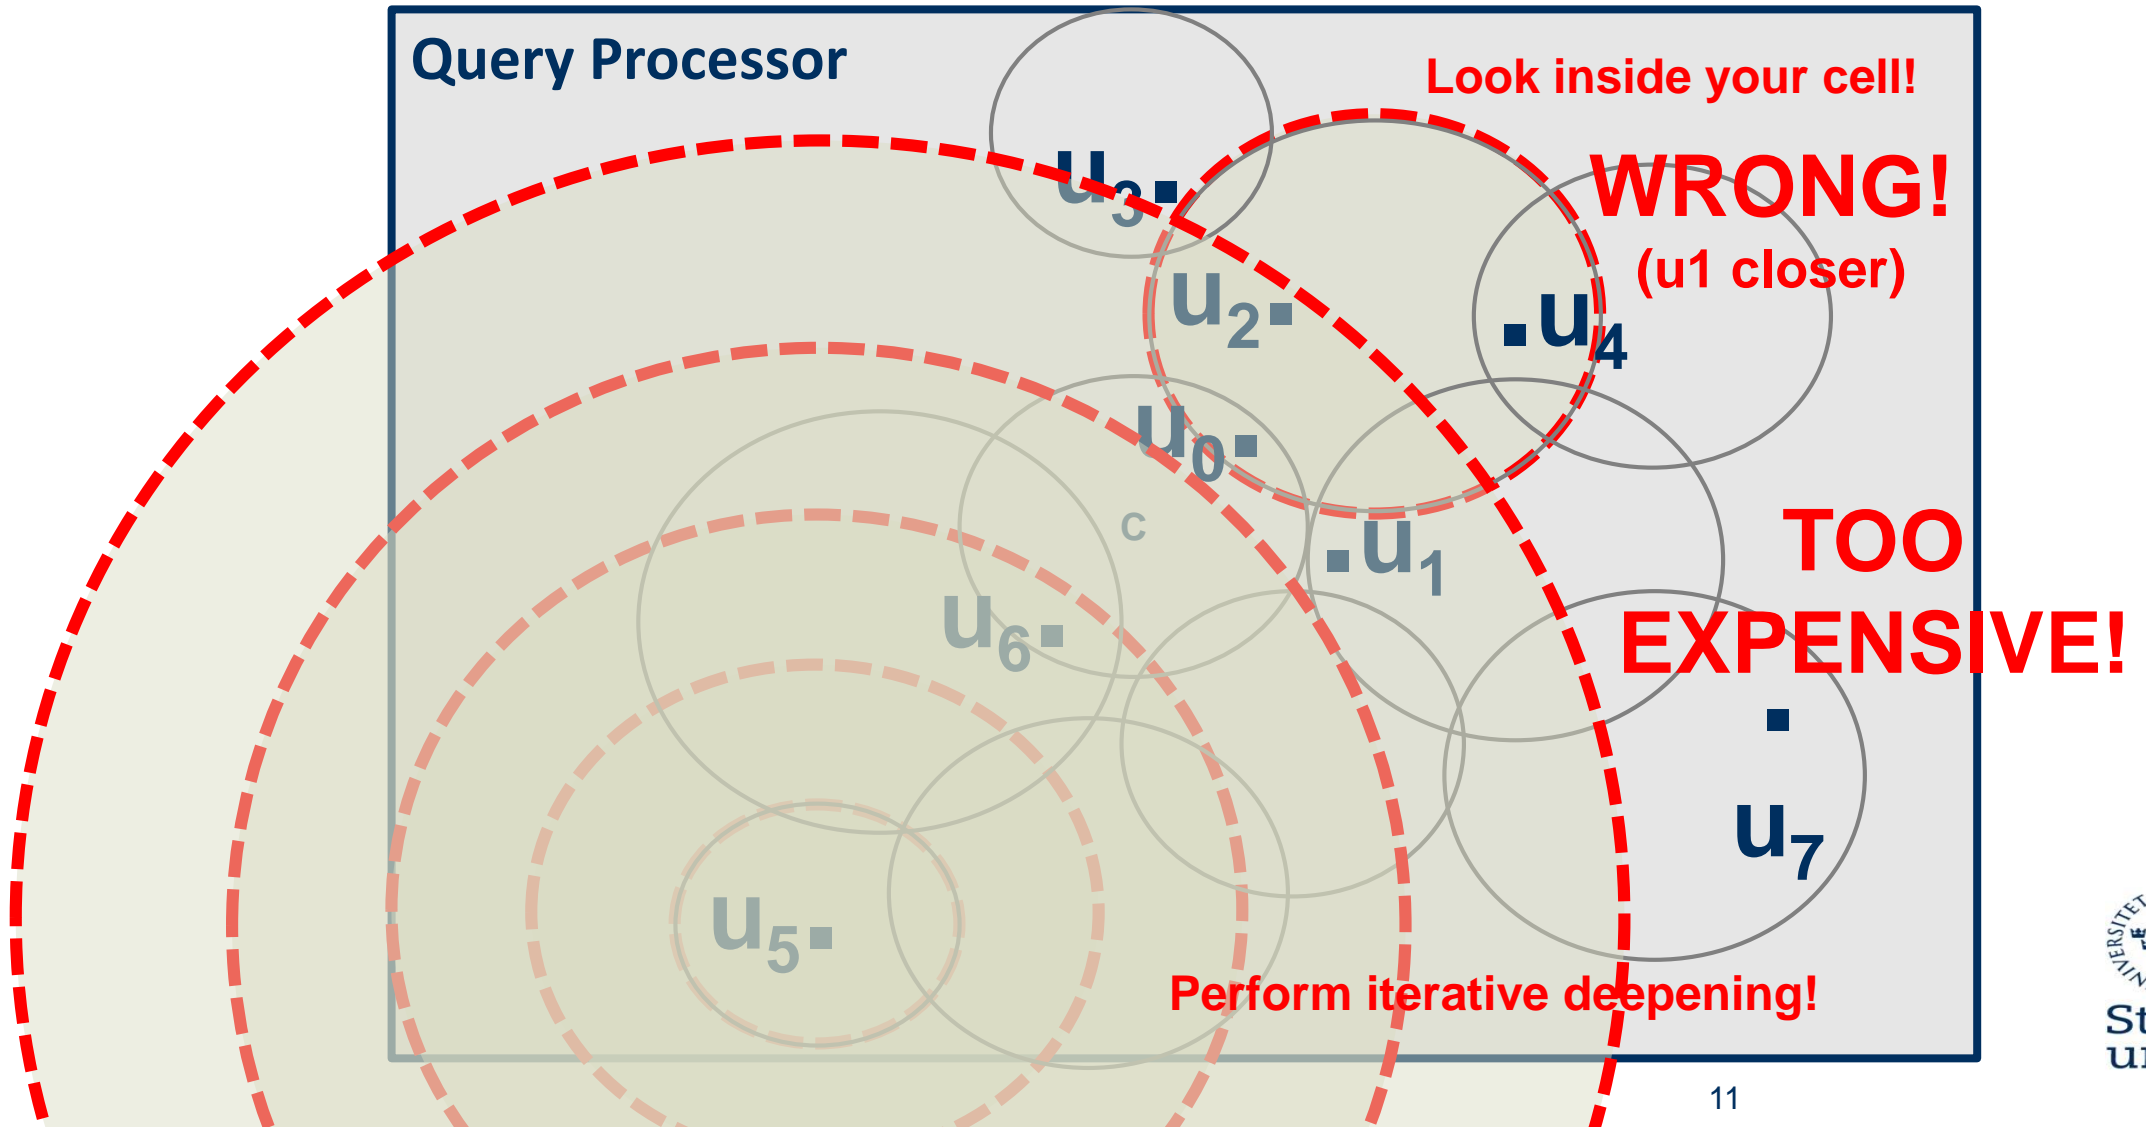
- **Definition of the CAkNN problem:**

Given a set U of n users and their location reports $r_{i,t} \in R$ at *timestep* $t \in T$, then the CAkNN problem is to find in each *timestep* $t \in T$ and for each user $u_i \in U$ the k objects $U_{sol} \subseteq U - u_i$ such that for all other objects $u_o \in U - U_{sol} - u_i$, $\text{dist}(u_k, u_i) \leq \text{dist}(u_o, u_i)$ holds



Naïve Solution

Find 2-NN for u_0 at timestep t . For u_5 ?



Stateless & Stateful

- Existing algorithms for **CkNN (not CAkNN)**
Yu et al. (YPK) ¹¹ and Mouratidis et al. (CPM) ¹²
- **Stateless Version:** *Iteratively expand the search space for each user into neighboring cells to find the kNN (like previous slide)*
- **Stateful Version:** Improve the stateless version by *utilizing previous state, under the following assumptions:*
 - i) Static Querying User (i.e., designated for point queries)
 - ii) Target users move slowly (i.e., state does not decay)
 - iii) Few Target users

11. Yu, Pu, Koudas. "Monitoring k-nearest neighbor queries over moving objects," ICDE '05

12. Mouratidis, Papadias, Hadjieleftheriou, "Conceptual partitioning: an efficient method for continuous nearest neighbor monitoring," SIGMOD '05.

Outline

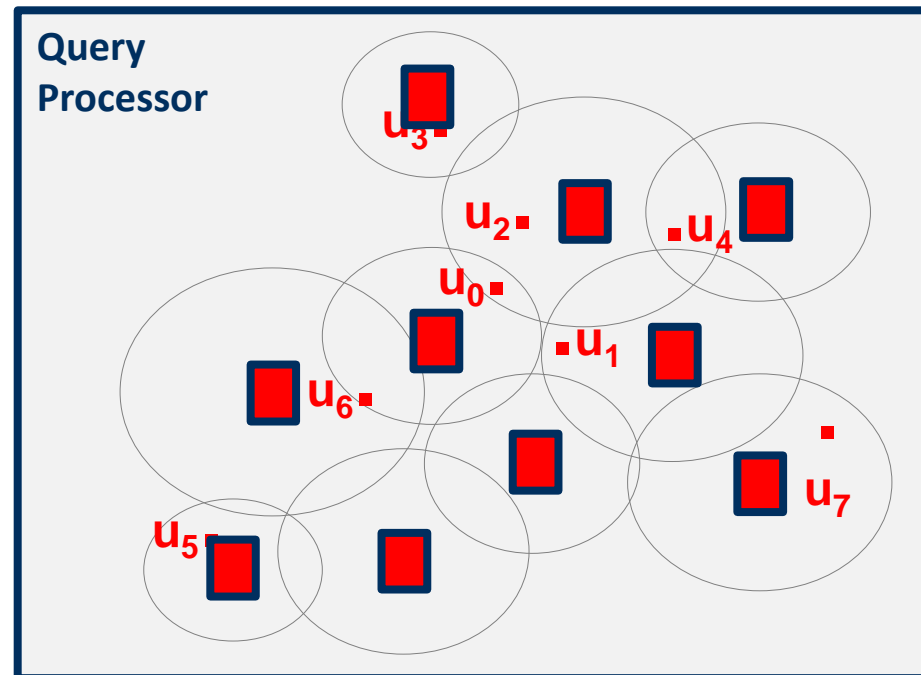
- Motivation
- System Model and Problem Formulation
- **Proximity Algorithm**
- Experimental Evaluation

Proximity Overview

- The first specialized algorithm for *Continuous All k-NN* (**CAkNN**) queries
- Important characteristics:
 - **Stateless** (i.e., optimized for **high mobility**)
 - **Batch processing** (i.e., with **search space searching**)
 - **Parameter-free** (i.e., no tuning parameters)
- Generic operator for proximity-based queries
 - See Crowdcast application presented later.

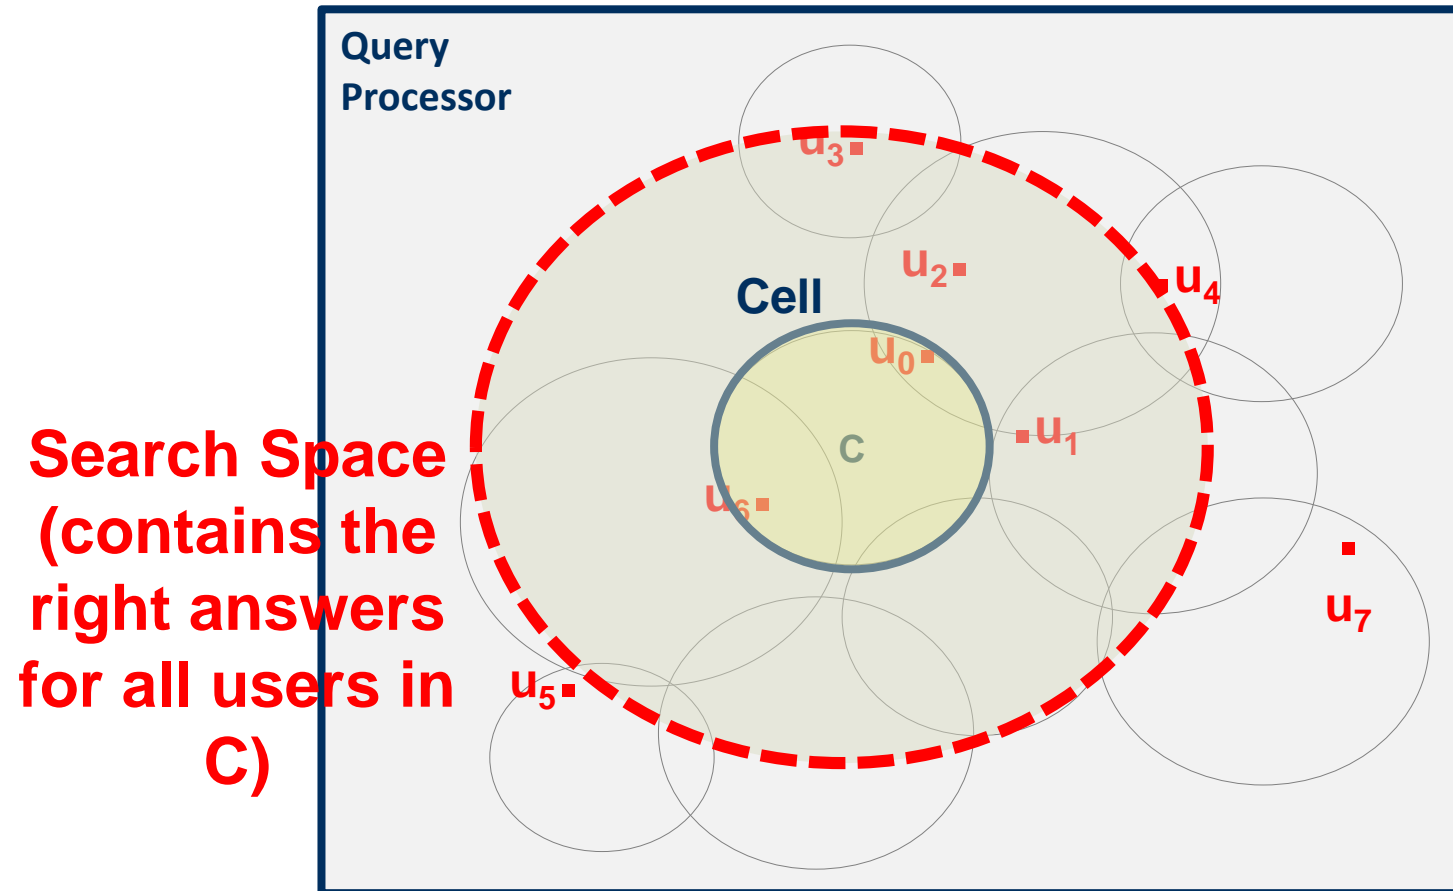
Proximity Outline

- For every timestep:
 1. Initialize a k^+ -heap for every cell
 2. Insert **every** user's location report to every k^+ -heap
 - Notice that k^+ -heap is a heap-based structure and most location reports will be dropped as a result of an insert operation
 3. For every user scan the k^+ -heap of his cell to find his k -NN



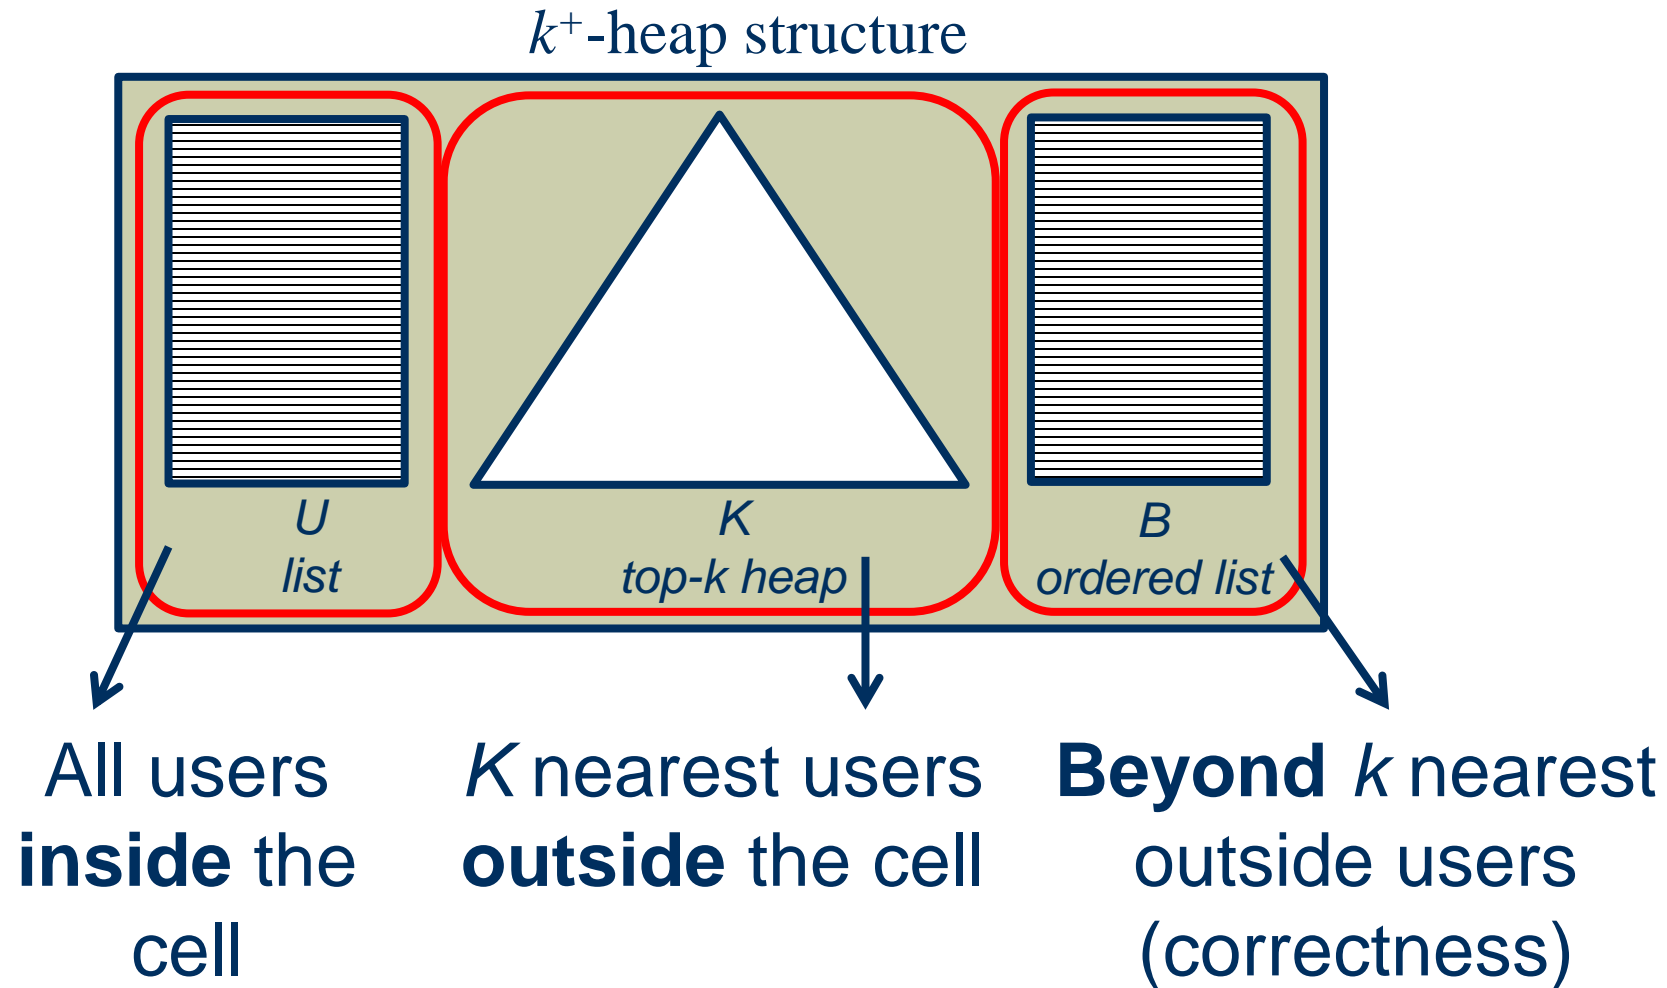
Intuition behind Proximity

- Users in a cell will share the **same search space** (search space sharing)
- Compute 1 search space per cell only!



Proximity k^+ -heap

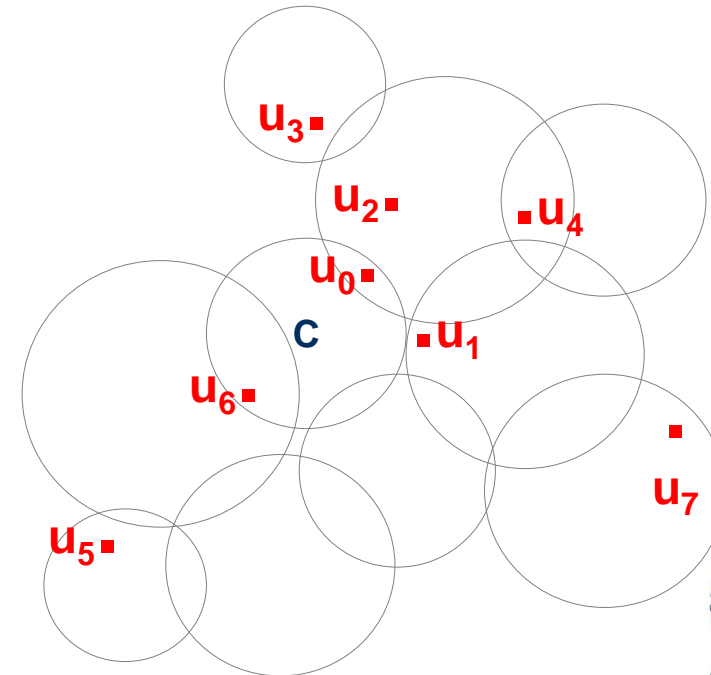
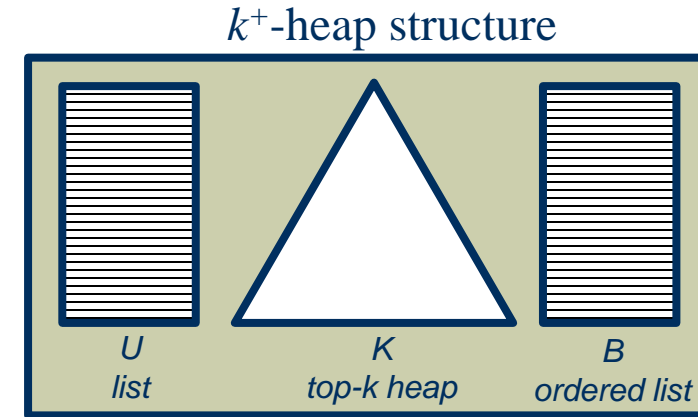
- *The k^+ -heap structure for a cell*



k^+ -heap Construction (for Cell C)

Assume $k=2$.

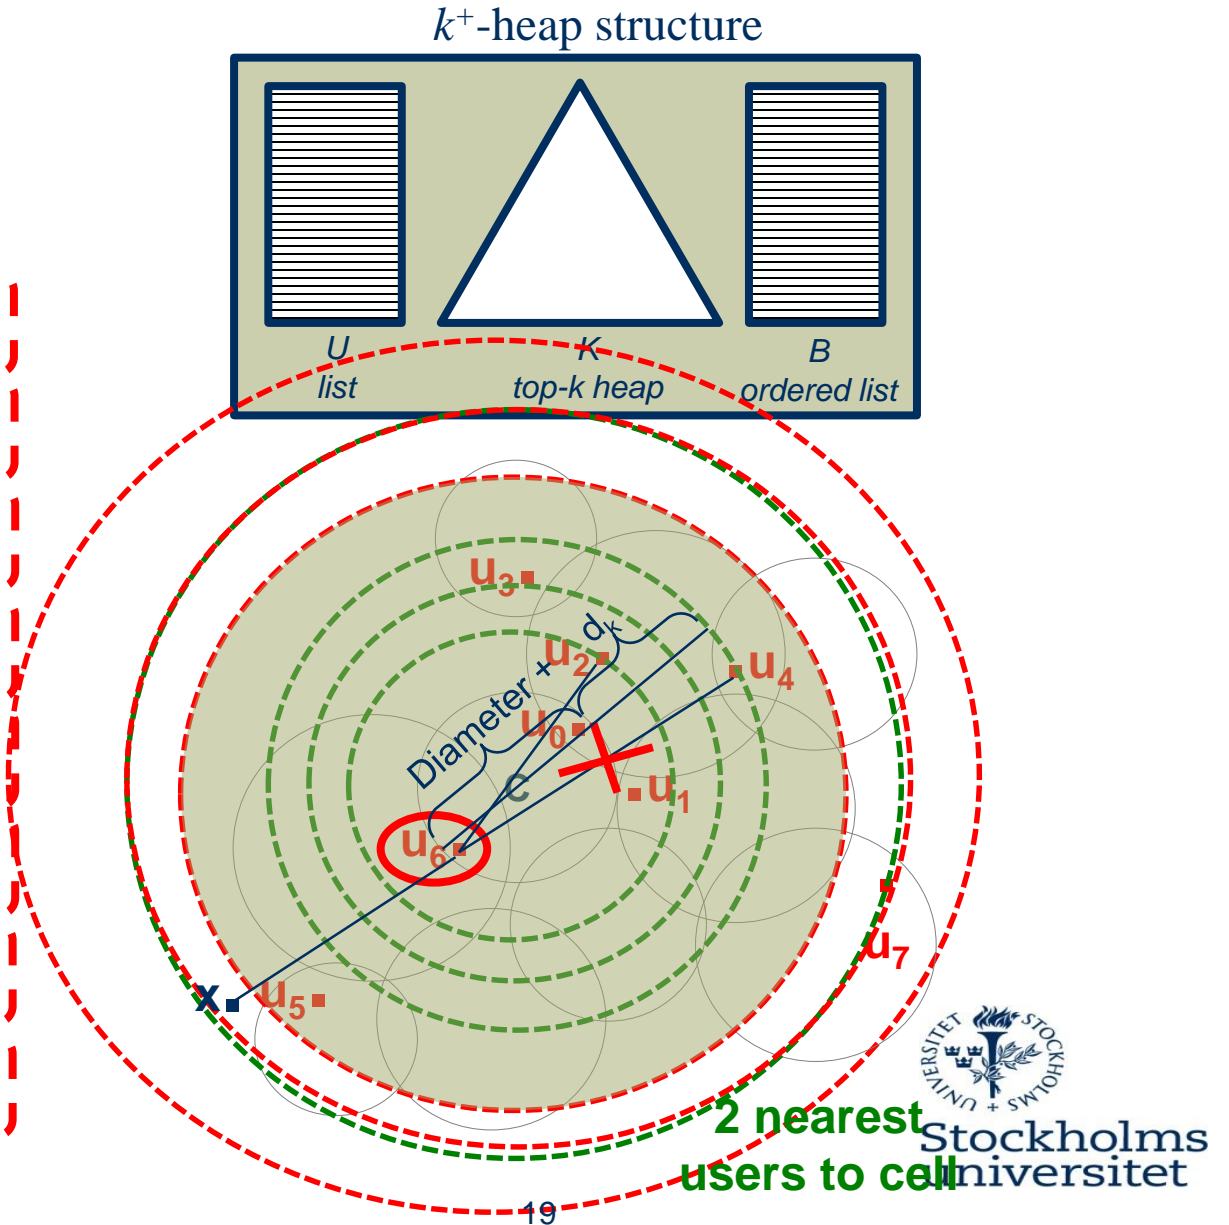
Arriving Reports	Structure U	Structure K	Structure B
u_6	u_6		
u_4	u_6	u_4	
u_7	u_6	u_7, u_4	
u_2	u_6	u_4, u_2	u_7
u_3	u_6	u_3, u_2	u_4, u_7
u_1	u_6	u_2, u_1	u_3, u_4
u_5	u_6	u_2, u_1	u_3, u_4, u_5
u_0	u_6, u_0	u_2, u_1	u_3, u_4, u_5



k^+ -heap Construction (for Cell C)

Assume $k=2$.

Arriving Reports	Structure U	Structure K	Structure B
u_6	u_6		
u_4	u_6	u_4	
u_7	u_6	u_7, u_4	
u_2	u_6	u_4, u_2	u_7
u_3	u_6	u_3, u_2	u_4, u_7
u_1	u_6	u_2, u_1	u_3, u_4
u_5	u_6	u_2, u_1	u_3, u_4, u_5
u_0	u_6, u_0	u_2, u_1	u_3, u_4, u_5



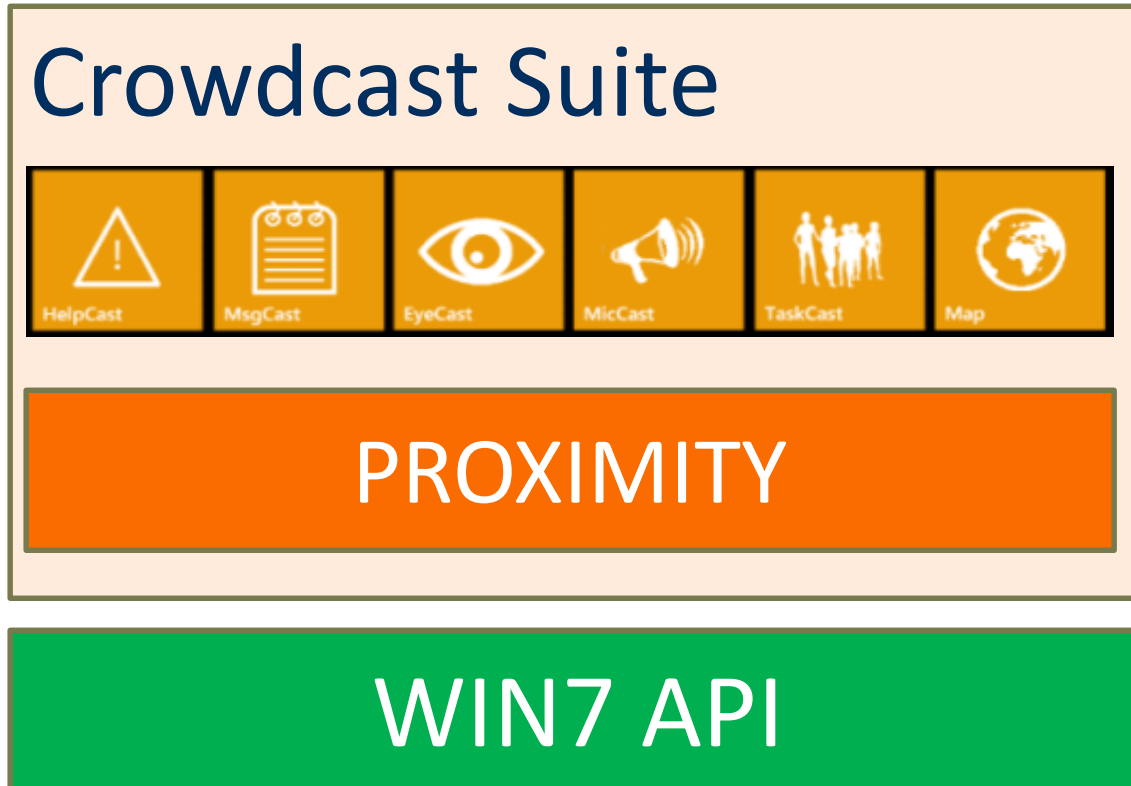
Outline

- Motivation
- System Model and Problem Formulation
- Proximity Algorithm
- **Experimental Evaluation**

Experimental Methodology

- **Dataset:**
 - Oldenburg Dataset:
 - Spatiotemporal generator with roadmap of Oldenburg as input (25km x 25km)
 - 5000 maximum users (Oldenburg population: 160.000)
 - Manhattan Dataset:
 - Vehicular mobility generator with roadmap of Manhattan as input (3km x 3km)
 - 500 users
- Network connectivity points (*NCPs*) uniformly in space
 - Communication ranges 1km, 4km and 16km for the *Oldenburg* dataset and ranges 1km and 4km for the *Manhattan* dataset
- **Query:** *CAkNN*
- **Comparison:** Proximity vs YPK vs CPM (using the optimal parameter value for YPK and CPM)

The Crowdcast Application Suite



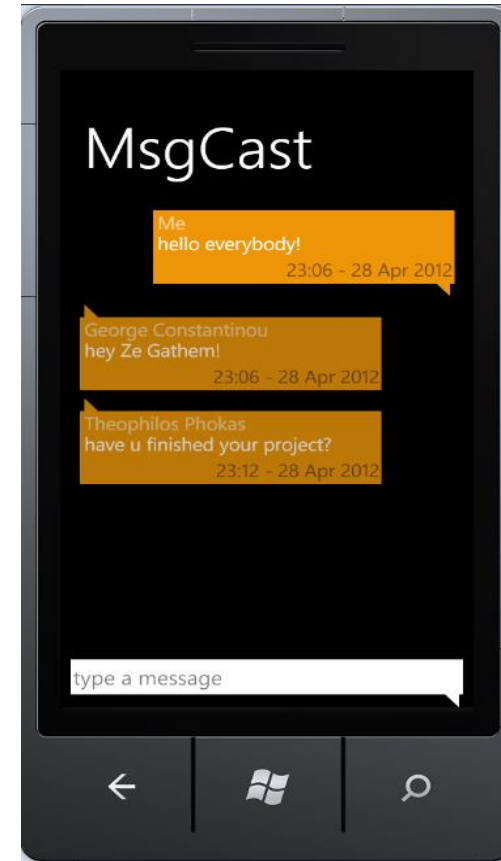
Ranked 3rd in Microsoft's *ImagineCup* contest local competition (to appear in the next demo session!)



Crowdcast: MsgCast



- Location-based Chat Channel (Post / Follow)
- Provide Guidelines

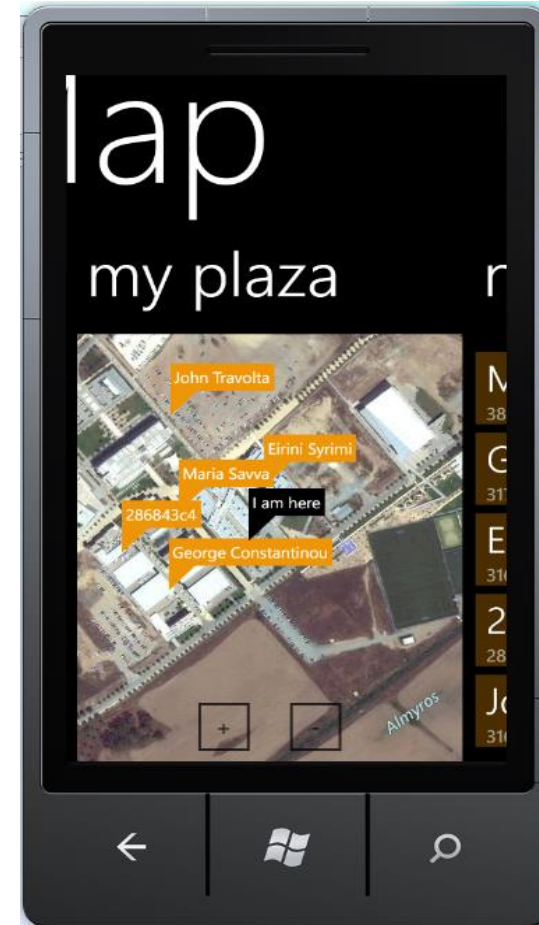
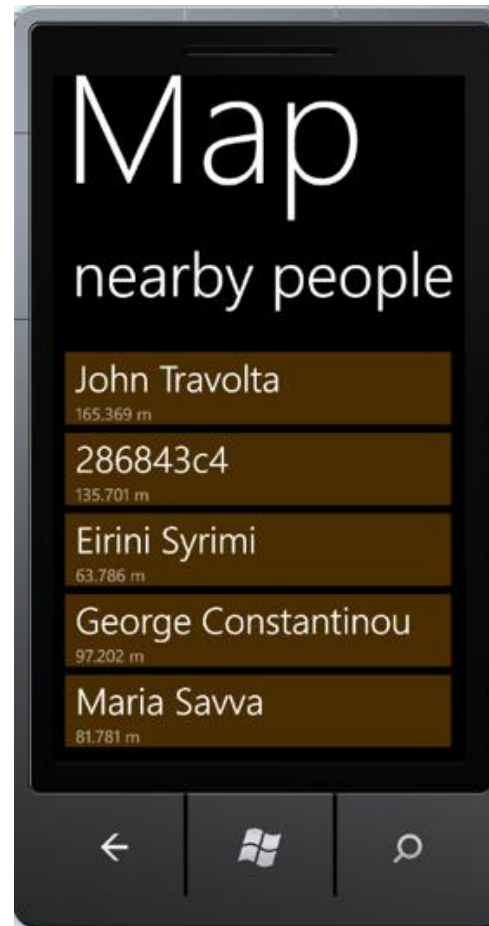


"Get your location-based questions answered!"

Crowdcast: HelpCast



***"The Ubiquitous
Help Platform for
Anyone in Need"***



Crowdcast: EyeCast

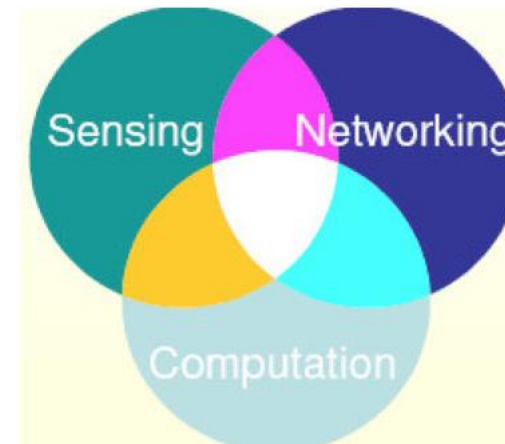


- Disaster Recovery Operations
- Indoor Video Conferencing Network



"Extend your Vision beyond your 2 eyes!"

IOT, Proximity Sensing Localization in Sensor Network Part II



Part II Outline

- GPS is not always good
- Localization Problem
- Many Ways to approach this problem
 - Trilateration and Triangulation
 - Ad-hoc approaches
 - Optimization
 - Multidimensional Scaling (MDS)
 - Fingerprinting, classification and scene analysis
- Multilateration: Use plane geometry
- Collaborative Multilateration: Use joint optimization
- Mass-spring Localization

Find Where the Sensor is

- Location information is important.
 1. Devices need to know where they are.
 - Sensor tasking: turn on the sensor near the window...
 2. We want to know where the data is about.
 - A sensor reading is too hot – where?
 3. It helps infrastructure establishment, such as geographical routing and sensor coverage.
 - Intruder detection;
 - Localized geographical routing.

GPS is not always good

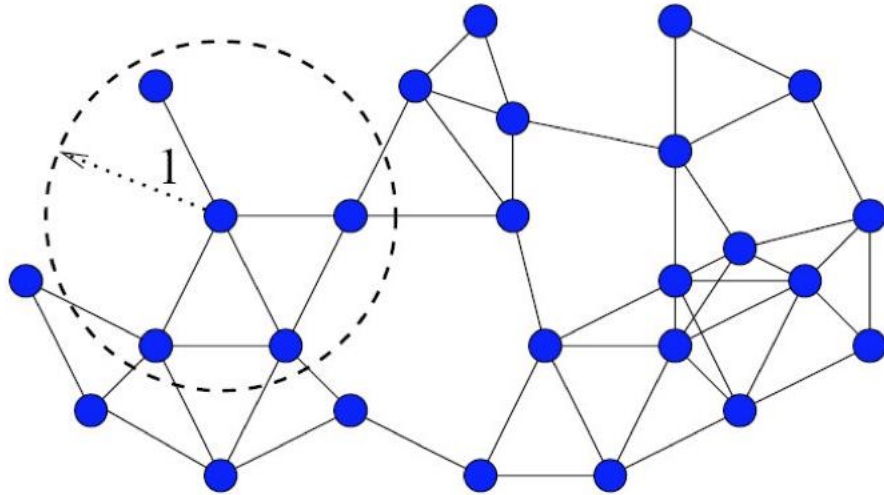
- Requires clear sky, doesn't work indoor.
- Too expensive.
 - A \$1 sensor attached with a \$100 GPS?

Localization:

- (optional) Some nodes (anchors or beacons) have GPS or know their locations.
- Nodes make local measurements;
 - Distances between two sensors, angles between two neighbors, etc.
- Communicate between each other;
- Infer location information from these measurements.

Model of a sensor network

- Sensor networks with omni-directional antennas are usually modeled by unit disk graphs.
 - Two nodes have a link if and only if their distance is within 1.



- Use the graph property (connectivity, local measurements) to deduct the locations.

Localization Problem

- Output: nodes' location.
 - Global location, e.g., what GPS gives.
 - Relative location.
- Input:
 - Connectivity, hop count.
 - Nodes with k hops away are within Euclidean distance k .
 - Nodes without a link must be at least distance 1 away.
 - Distance measurement of an incoming link.
 - Angle measurement of an incoming link.
 - Combinations of the above.

Measurements

Distance estimation:

- Received Signal Strength Indicator (RSSI)
 - The further away, the weaker the received signal.
 - Mainly used for RF signals.
- Time of Arrival (ToA) or Time Difference of Arrival (TDoA)
 - Signal propagation time translates to distance.
 - RF, acoustic, infrared and ultrasound.

Angle estimation:

- Angle of Arrival (AoA)
 - Determining the direction of propagation of a radio-frequency wave incident on an antenna array.
- Directional Antenna
- Special hardware, e.g., laser transmitter and receivers.

Localization

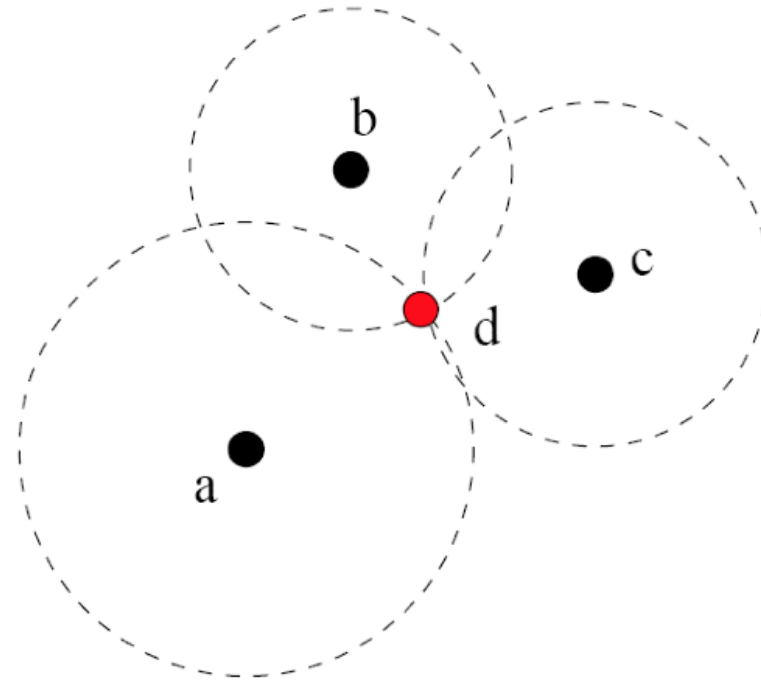
- Given distances or angle measurements, find the locations of the sensors.
- **Anchor-based**
 - Some nodes know their locations, either by a GPS or as pre-specified.
- **Anchor-free**
 - Relative location only.
 - A harder problem, need to solve the global structure. Nowhere to start.
- **Range-based**
 - Use range information (distance estimation).
- **Range-free**
 - No distance estimation, use connectivity information such as hop count.

Many Ways to approach this problem

- Trilateration and triangulation
- Fingerprinting and classification
- Ad-hoc and range/free
- Graph rigidity
- Identifying codes
- Bayesian Networks
- Optimization
- Multi-dimensional scaling

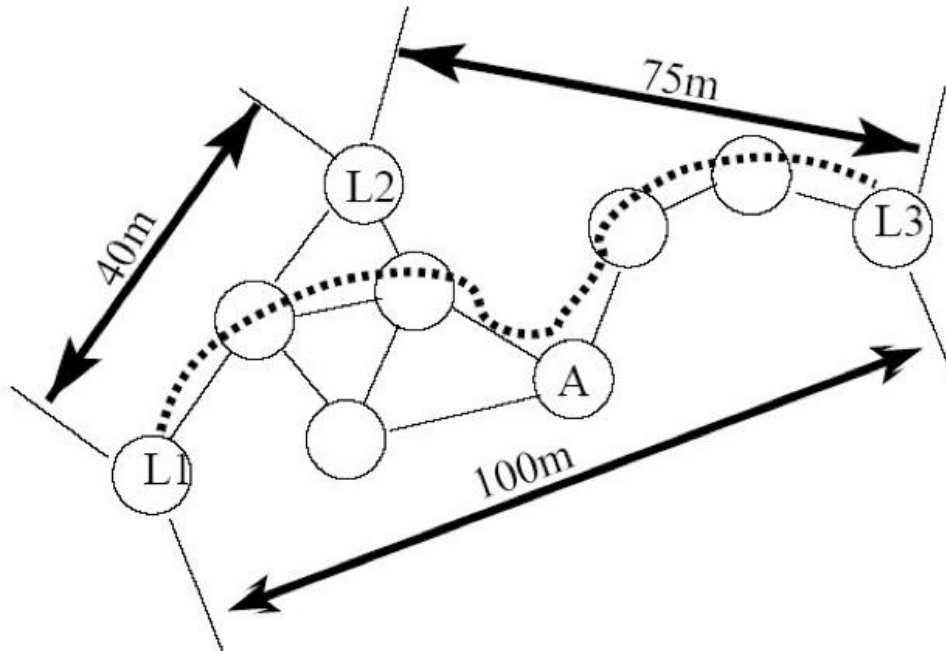
Trilateration and Triangulation

- Use geometry, measure the distances/angles to three anchors.
- Trilateration: use distances
 - Global Positioning System (GPS)
- Triangulation: use angles
 - Cell phone systems.
- How to deal with inaccurate measurements?
- How to solve for more than 3 (inaccurate) measurements?



Ad-hoc approaches

- Ad-hoc positioning (APS)
 - Estimate range to landmarks using hop count or distance summaries
- APS:
 - Count hops between landmarks
 - Find average distance per hop
 - Use multi-lateration to compute location



Optimization

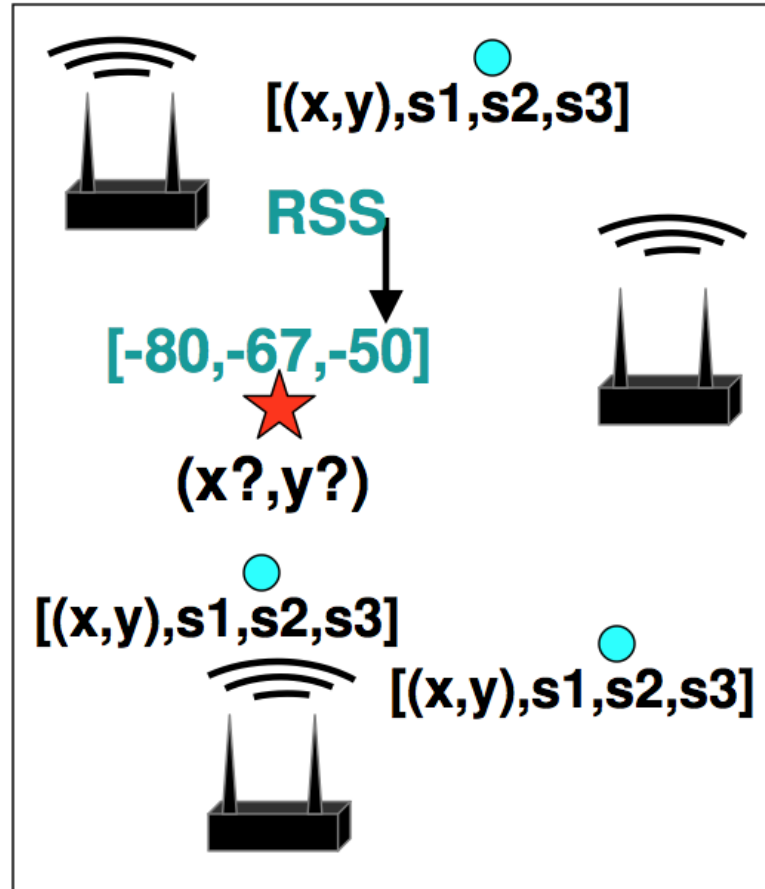
- View system of nodes, distances and angles as a system of equation with unknowns.
- Add inequalities
 - E.g. radio range is at most 1.
- Solve resulting system of inequalities as an optimization problem.

Multidimensional Scaling (MDS)

- MDS is a general method of finding points in a geometric space that preserves the pair-wise distances as much as possible.
 - It works also for non-metric data.
- Given the pairwise distances P , find a set of points X in m -dimensional space.
- Take the largest 2 eigenvalues and eigenvectors of X as the best 2D approximations.

Fingerprinting, classification and scene analysis

- Offline phase: collect training data (fingerprints): $[(x, y), SS]$.
 - E.g., the mean Signal Strength to N landmarks.
- Online phase: Match RSS to existing fingerprints probabilistically or by using a distance metric.
- Cons:
 - How to build the map?
 - Someone walks around and samples?
 - Automatic?
 - Sampling rate?
 - Changes in the scene (people moving around in a building) affect the signal strengths.



Bayesian Networks

- View positions as random variables
- Build network to describe likely values of these variables given observations
- Pros:
 - Captures any set of observations and priors
- Cons:
 - Computationally expensive
 - Accuracy

Reference Papers

- **Multi-lateration:**
- **[Savvides01]** A. Savvides, C.-C. Han, and M. B. Strivastava. Dynamic fine-grained localization in ad-hoc networks of sensors. Proc. MobiCom 2001.
- **[Savvides03]** A. Savvides, H. Park, and M. B. Strivastava. The n -hop multilateration primitive for node localization problems, Mobile Networks and Applications, Volume 8, Issue 4, 443-451, 2003.
- **Mass-spring model:**
- **[Howard01]** A. Howard, M. J. Mataric, and G. Sukhatme, Relaxation on a Mesh: a Formalism for Generalized Localization, IEEE/RSJ International Conference on Intelligent Robots and Systems, October, 2001.